

# Workshop: Data science with R (ZEW)

## Session #7: Supervised Learning

Obryan Poyser  
2019-04-18



# Outline

- What is Supervised Learning?
- Machine Learning vs Econometrics
  - Prediction/classification vs inference
- Parametrics vs Non-parametric estimation
- Classification
  - Logistic (softmax)
  - Performance measures
- Prediction
  - Regularization, shrinkage and variable selection
  - Sparsity, Ridge, and LASSO regressions
  - Ensemble learning
  - Performance measures

# What is supervised learning?

- Broadly speaking the main two subfields of machine learning are supervised learning and unsupervised learning.
- Supervised statistical involves building statistical models  $f()$  for **predicting** or **estimating** and *output* ( $Y$ ) based on one or more *inputs* given by the design matrix ( $X$ ).

$$\hat{Y} = \hat{f}(X)$$



# Why estimate $f$ ?

## Prediction & classification

- In classification, the goal is to predict a *class label* within a defined set of elements.
- Prediction is mostly associated with continuous data.
- $\hat{f}$  could be treated as a "*black box*", that is, we are not concerned on the form of  $\hat{f}$ , instead, how good this function predicts  $\hat{Y}$ .
- No matter how accurate  $\hat{f}$  is (by choosing a statistical learning technique), there always be non-reducible error term (irreducible error).
- We can improve  $\hat{f}$  by choosing a more flexible form. But it has a cost!

## Inference


- We want to understand the effect of  $X_1, \dots, X_p$  on  $Y$ .
- By definition, it can not be a "*black box*" function, because we need to know the exact form.
- ML statistics differentiate from Econometrics in this element.
  - Econometrics: "*the quantitative analysis of actual economic phenomena based on the concurrent development of theory and observation, related by appropriate methods of inference.*" [source](#)
  - Machine learning: "*Machine Learning is the field of study that gives computers the ability to learn without being explicitly programmed*" Arthur Samuel, 1959.
- Causal inference is getting attention from ML supporters, but there is much work to do.

# Econometrics vs Machine Learning: the $\hat{\beta}$ vs $\hat{y}$ dilemma

- Many economic applications revolve around *parameter estimation*
  - Produce good estimates that unveil the true relationship between  $y$  and  $X$
- Machine learning algorithms are not designed for inference purposes
  - One has to be aware of the properties and goals of the estimators, the typical parameters' interpretation (i.e. asymptotic theory least square) is no necessary longer valid.
- Applications
  - New data for traditional questions: for instance: measuring the level of economic activity from satellite maps using light-intensity measures.
  - Pre-processing
    - Propensity Score Matching, Linear Instrumental Variables Regression, Heterogeneous treatment effects.
- ML algorithms are technically easy to use in Python or R
  - Threats: *naive interpretations!*




# Predicting poverty from satellite maps



# How to estimate $\hat{f}$ ?

## Parametric methods

- There is a reason for imposing a function form of  $f$ : interpretability. In econometrics the conditional expectation function  $E(Y|X = x) = f(x)$ , we restrict  $f$  from a theoretical grounds. Nonetheless, the CEF is by default nonparametric.
- Generalized Linear Model (GLMs) assume the design matrix can be expressed as a linear combination of a set of parameters (typically denoted as  $\beta$ ) obtained by OLS. The estimated parameters are linear by construction, while the  $X = x_i$  can be of any form or distribution, however, some transformation may be captured easier. Example:  
$$\text{income} = \beta_0 + \beta_1 \text{education} + \beta_2 \text{seniority}$$

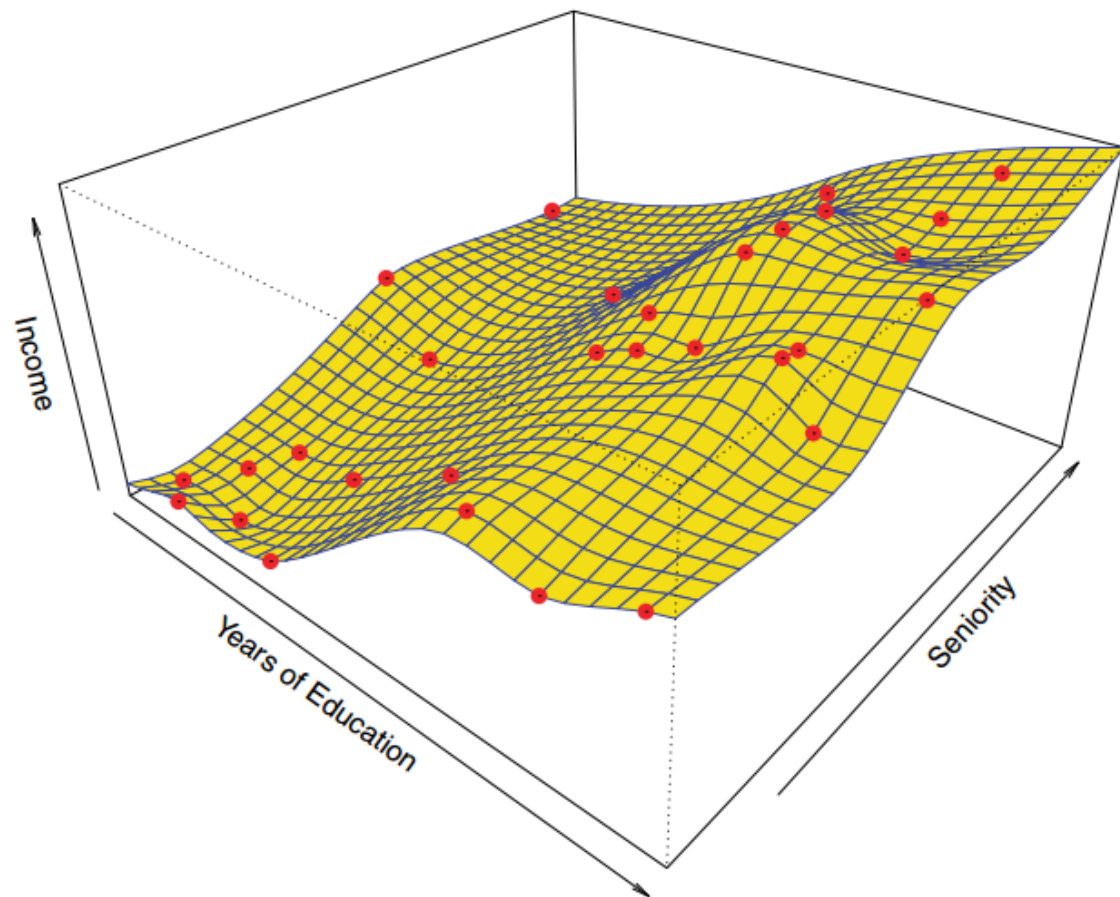


Parametric Linear Model. Hastie et al. 2013.

# How to estimate $\hat{f}$ ?

## Non-parametric methods

- Non-parametric methods do not make explicit assumptions about the functional form of  $f$ .
- Avoiding the assumption of a particular functional form for  $f$ , they have the potential to accurately fit a wider range of possible shapes for  $f$ .
- Disadvantage: the amount of parameters increases drastically!
- Examples:
  - Multiple Adaptive Regression Splines:
$$\hat{f}(X) = \sum_{j=1}^k c_i B_i(X_j)$$
  - Generalized Additive Models:
$$\hat{f}(X) = s_0 + \sum_{j=1}^p s_j(X_j)$$
  - Random Forest
  - Neural Network
  - ...



Non-parametric model. Hastie et al. 2013.




# Trade-offs of selecting $f$

## Interpretation wise

- The more complex (flexible) the function form, the best it will fit the design matrix data generating process.
- The simplest a functional form is, the more interpretable are their parameters (GLM for instance), but, its fit is generally worse.

## Generalization wise

- Complex models tend to do a good job explaining the data used to estimate  $f$ , but it does a very bad job explaining data not used in the estimation process, this is called **overfitting**.
- Complex models often have **hyperparameters**, also name smoothing or complexity parameters, that cannot be estimated from the data. Examples: penalty term, the width of the kernel.
- **Underfitting** happens when the model is too simple,



Müller, Guido (2017)

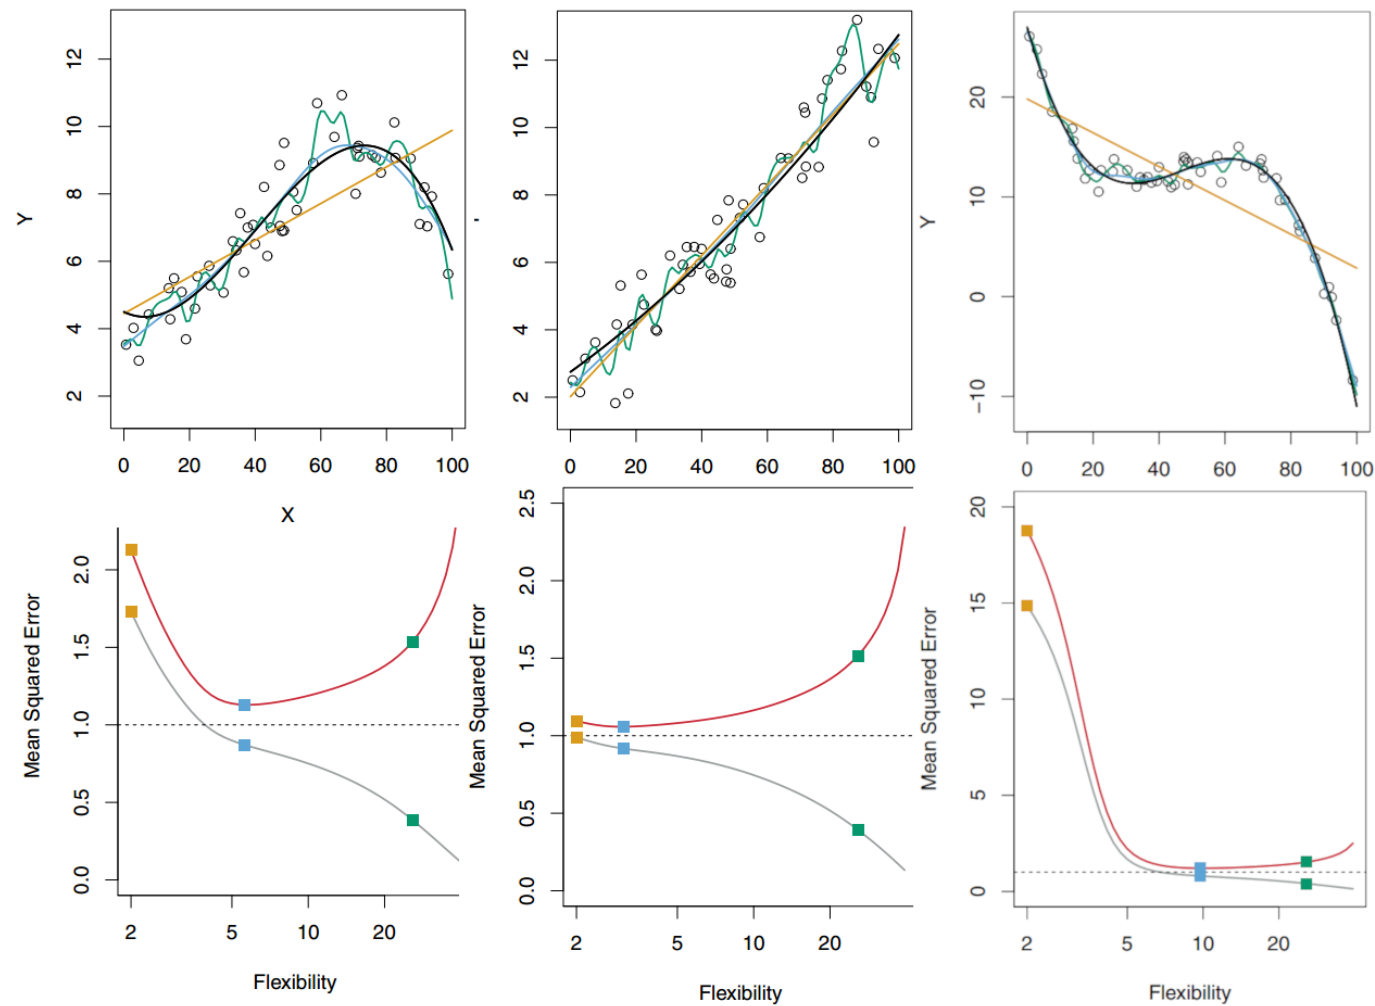
# Accuracy

- In order to evaluate the performance of a statistical learning method on a given data set, we need some way to measure how well its predictions actually match the observed data.
- We are interested in the accuracy of the *unseen* data, that is, a subsample of information which has not been used to estimate  $\hat{f}$ .
  - In practice, one has to divide the total number of observation into **training** and **testing** subsets. The first step is to use the **train** data and observe how good it predicts the **test** data.
  - As mentioned, complex models do a good job finding a function that captures as much variance for training design matrix, but they tend to extrapolate badly on **new information**, which make them useless.
- The fundamental problem of selecting  $\hat{f}$  is the trade-off between **bias** and **variance**.

$$E(Y - \hat{f}(X))^2 = Var(\hat{f}(X)) + [Bias(\hat{f}(X))]^2 + Var(\epsilon)$$


- Variance refers to the amount by which  $\hat{f}$  would change if we estimated it using a different training data set.
- Bias refers to the error that is introduced by approximating a real-life problem, which may be extremely complicated, by a much simpler model.

# Bias variance trade-off



# Lasso and Ridge regression

- Least Square estimates present commonly two inconvenient
  - Low prediction accuracy: low bias but large variance.
  - Interpretation: With a lower set of predictors is easier to visualize the "*big picture*".
  - Subset selection can be applied following *(for)backward-stepwise selection ad-hoc methods*.
- Another alternative is to use shrinkage (regularization) techniques that penalized parameter space, and reduce the sparsity of the design matrix, that is, when  $N < X$ .
- The goal is to reduce the space of regressors by imposing restrictions. Examples:
  - Spike and slab: In Bayesian models, it is possible to impose a strong prior to shrinkage the parameter space.
  - Principal component regression
  - **Ridge and Lasso**



Spike and Slab priors. (Poyser, 2018)

## Ridge regression

- Shrinks the regression coefficients by imposing a penalty on their size
- The ridge coefficients minimize a penalized residual sum of squares

$$\begin{aligned}\hat{\beta}^{ridge} &= \arg \max_{\beta} \left\{ \sum_{i=1}^N \left( y_i - \beta_0 - \sum_{j=1}^p x_{ij} \beta_j \right)^2 + \lambda \sum_{j=1}^p \beta_j^2 \right\} \\ &= (X^T X + \lambda I)^{-1} X^T y\end{aligned}$$

Here  $\lambda \geq 0$  is the tuning parameter and controls the complexity of the estimation. The greater  $\lambda$  the greater the shrinkage.

- As expected, scaling affects the shrinkage component, therefore is necessary to standardize the inputs.
  - **Coefficients are no longer scale invariant**
- Ridge uses the so-called  $L_2$  norm

## Lasso regression

- Similar to Ridge, but instead of shrinking the coefficients with a squared rule of the parameters, it imposes an absolute value
- The penalty is a  $L_1$  norm  $||$
- Solution in  $y_i$  is non-linear
- Lasso does a continuous subset selection


$$\hat{\beta}^{lasso} = \arg \max_{\beta} \left\{ \sum_{i=1}^N \left( y_i - \beta_0 - \sum_{j=1}^p x_{ij} \beta_j \right)^2 + \lambda \sum_{j=1}^p |\beta_j| \right\}$$

- Lasso  $\lambda$  shrinkage paramets is chosen with the goal of minimizen the expected prediction error.

# Decision-tree based models

- Decision trees are widely used models for classification and regression tasks.
- Essentially, they learn a hierarchy of if/else questions, leading to a decision.
- Tree-based methods are simple and useful for interpretation. However, they typically are not competitive with the best-supervised learning approaches
- $R_1, R_2, R_3$  are known as *leaves* or *terminal nodes* of the tree
- Steps:
  - Dividing the predictor space  $X_p$  into  $R_1, \dots, R_J$  non-overlapping regions
  - For every observation that falls into  $R_J$ , take the mean of the dependent variable
- The objective is to minimize the Residual Sums of Squares of:

$$\sum_{j=1}^J \sum_{i \in R_j} (y_i - \hat{y}_{R_j})^2$$




# Bagging

- Decision trees have the advantage of interpretability, however, they tend to show high variance
- One way to reduce the variance is to take the average of all the predictions
- Since the actual design matrix does not change, it is not possible to take the average of inexistent samples. To solve this problem we can use **bootstrapping**
  - What bootstrap does is taking  $B$  random samples with replacement from a single sample, creating several artificial samples?
- Bagging uses a bootstrap rule to create pseudo-training sets, then average the prediction in order to reduce the variance.

$$\hat{f}_{bag}(x) = \frac{1}{B} \sum_{b=1}^B \hat{f}^{*b}(x)$$

- Even though bagging improves accuracy, it reduces interpretability. Nevertheless, there is one way to assess the importance of one variable: variable importance
- The process works as follows:
  - Take one variable at a time, then calculate how much the RSS decreases in each split

# Bagging



Bagging process. Géron (2017)



# Random-forest

- Works similar to bagging, but with the difference of taking a random sample of predictors over the design matrix  $X_p$
- Process:
  - Draw a bootstrap sample  $Z^*$  of size  $N$  from the training data.
  - Grow a random-forest tree  $T_b$  to the bootstrapped data, by recursively repeating the following steps for each terminal node of the tree, until the minimum node size  $n_{min}$  is reached.
  - Select  $m$  variables at random from the  $p$  variables.
  - Pick the best variable/split-point among the  $m$ .
  - Split the node into two daughter nodes.

## Considerations

- There are several algorithms and tweaks to improve prediction accuracy, however, each of them deserves a whole session to be studied.
  - Boosting
  - Ensemble
  - Boosting
  - SVM
  - Neural Networks
  - ...
- For a more detailed explanation of such techniques see the [references](#) section.

# Packages

- `broom`: is an attempt to bridge the gap from untidy outputs of predictions and estimations to the tidy data we want to work with.
  - `tidy`: constructs a data frame that summarizes the model's statistical findings. This includes coefficients and p-values for each term in a regression, per-cluster information in clustering applications, or pre-test information for multtest functions.
  - `augment`: add columns to the original data that was modeled. This includes predictions, residuals, and cluster assignments.
  - `glance`: construct a concise one-row summary of the model. This typically contains values such as  $R^2$ , adjusted  $R^2$ , and residual standard error that are computed once for the entire model.
- `parsnip`
  - It is designed to solve a specific problem related to model fitting in R, **the interface**.
  - Many functions have different interfaces and arguments names and parsnip standardizes the interface for fitting models as well as the return values
- `tidymodels`: Collection of modeling packages, that aim to create a common structure (similar to sklearn in Python)
- `rsample`: Classes and functions to create and summarize different types of resampling objects (e.g. bootstrap, cross-validation).

# Dataset

## Wine quality

```
wine_w <- readr::read_rds("datasets/session_7/wine_white.rds")
skimr::skim(wine_w) %>% skimr::kable()
```

## Skim summary statistics














## n obs: 4898

## n variables: 13

##

## Variable type: numeric

##

##	variable	missing	complete	n	mean	sd	p0	p25	p50	p75	p100	hist
##	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----
##	alcohol	0	4898	4898	10.51	1.23	8	9.5	10.4	11.4	14.2	
##	chlorides	0	4898	4898	0.046	0.022	0.009	0.036	0.043	0.05	0.35	
##	citric_acid	0	4898	4898	0.33	0.12	0	0.27	0.32	0.39	1.66	
##	density	0	4898	4898	0.99	0.003	0.99	0.99	0.99	1	1.04	
##	fixed_acidity	0	4898	4898	6.85	0.84	3.8	6.3	6.8	7.3	14.2	
##	free_sulfur_dioxide	0	4898	4898	35.31	17.01	2	23	34	46	289	
##	p_h	0	4898	4898	3.19	0.15	2.72	3.09	3.18	3.28	3.82	
##	quality	0	4898	4898	5.88	0.89	3	5	6	6	9	
##	quality_b	0	4898	4898	0.67	0.47	0	0	1	1	1	
##	residual_sugar	0	4898	4898	6.39	5.07	0.6	1.7	5.2	9.9	65.8	
##	sulphates	0	4898	4898	0.49	0.11	0.22	0.41	0.47	0.55	1.08	
##	total_sulfur_dioxide	0	4898	4898	138.36	42.5	9	108	134	167	440	
##	volatile_acidity	0	4898	4898	0.28	0.1	0.08	0.21	0.26	0.32	1.1	

## Training, testing



# Cross-validation

Split 1	Fold 1	Fold 2	Fold 3	Fold 4	Fold 5	Metric 1
Split 2	Fold 1	Fold 2	Fold 3	Fold 4	Fold 5	Metric 2
Split 3	Fold 1	Fold 2	Fold 3	Fold 4	Fold 5	Metric 3
Split 4	Fold 1	Fold 2	Fold 3	Fold 4	Fold 5	Metric 4
Split 5	Fold 1	Fold 2	Fold 3	Fold 4	Fold 5	Metric 5

Training	Testing
----------	---------

# Training, testing, and cross-validation

```
data_split <- initial_split(wine_w, prop = .7)  
(train_data <- training(data_split)) %>% dim()
```

```
## [1] 3429  13
```

```
(test_data <- testing(data_split)) %>% dim()
```

```
## [1] 1469  13
```

# Cross-validation

```
(cv_data <- vfold_cv(train_data, v = 10) %>%  
  mutate(train=map(splits, ~training(.x))  
    , validate=map(splits, ~testing(.x))  
    , lm=map(train, ~lm(quality~.-quality_b, data = .x))  
    , rf=map(train, ~rand_forest(mtry = 5, trees = 200) %>%  
      set_engine("ranger", importance="impurity") %>%  
      fit(quality~.-quality_b, data=.x)  
    )  
  )  
)
```

## # 10-fold cross-validation

## # A tibble: 10 x 6

##	splits	id	train	validate	lm	rf
##	* <list>	<chr>	<list>	<list>	<list>	<list>
##	1 <split [3.1K/3...	Fold01	<tibble [3,086 x...	<tibble [343 x...	<S3: l...	<fit[+...
##	2 <split [3.1K/3...	Fold02	<tibble [3,086 x...	<tibble [343 x...	<S3: l...	<fit[+...
##	3 <split [3.1K/3...	Fold03	<tibble [3,086 x...	<tibble [343 x...	<S3: l...	<fit[+...
##	4 <split [3.1K/3...	Fold04	<tibble [3,086 x...	<tibble [343 x...	<S3: l...	<fit[+...
##	5 <split [3.1K/3...	Fold05	<tibble [3,086 x...	<tibble [343 x...	<S3: l...	<fit[+...
##	6 <split [3.1K/3...	Fold06	<tibble [3,086 x...	<tibble [343 x...	<S3: l...	<fit[+...
##	7 <split [3.1K/3...	Fold07	<tibble [3,086 x...	<tibble [343 x...	<S3: l...	<fit[+...
##	8 <split [3.1K/3...	Fold08	<tibble [3,086 x...	<tibble [343 x...	<S3: l...	<fit[+...
##	9 <split [3.1K/3...	Fold09	<tibble [3,086 x...	<tibble [343 x...	<S3: l...	<fit[+...
##	10 <split [3.1K/3...	Fold10	<tibble [3,087 x...	<tibble [342 x...	<S3: l...	<fit[+...

# Performance

## Linear model

```
lm_model <- cv_data %>%
  mutate(pred=map2(.x=lm, .y=validate, ~predict(.x, .y) %>%
    as_tibble %>%
    set_names("pred")))
  , mae=map2_dbl(.x=pred, .y=validate
    , ~measures::MAE(truth = .y$quality
    , response = .x$pred
    )
  , mse=map2_dbl(.x=pred, .y=validate
    , ~measures::MSE(truth = .y$quality
    , response = .x$pred
    ))
```


## Random forest

```
rf_model <- cv_data %>%
  mutate(pred=map2(.x=rf, .y=validate, ~predict(.x, .y) %>%
    as_tibble %>%
    set_names("pred")))
  , mae=map2_dbl(.x=pred, .y=validate
    , ~measures::MAE(truth = .y$quality
    , response = .x$pred
    )
  , mse=map2_dbl(.x=pred, .y=validate
    , ~measures::MSE(truth = .y$quality
    , response = .x$pred
    ))
```



# Model performance

```
lm_model %>%
  dplyr::select(id, mae, mse) %>%
  bind_rows(rf_model %>%
    dplyr::select(id, mae, mse)
    , .id = "model") %>%
  mutate(model=case_when(
    model=="linear" ~ "linear"
    , T ~ "random_forest"
  )) %>%
  pivot_longer(cols = c(mae, mse)
    , names_to = "measure") %>%
  ggplot(aes(model, value))+
  geom_boxplot() +
  geom_jitter(aes(col=id))+
  facet_grid(~measure, scales = "free")
```



# Hyperparameter tuning: grid search

```
(rf_tun <- vfold_cv(train_data, v = 5) %>%  
  crossing(mtry=1:5) %>%  
  mutate(train=map(splits, ~training(.x))  
    , validate=map(splits, ~testing(.x))  
    , rf=map2(.x=train, .y=mtry, ~rand_forest(mtry = .y, trees = 200) %>%  
      set_engine("ranger", importance="impurity") %>%  
      fit(quality~.-quality_b, data=.x)  
    ))  
)
```

## # A tibble: 25 x 6


##	splits	id	mtry	train	validate	rf
##	<list>	<chr>	<int>	<list>	<list>	<list>
##	1 <split [2.7K/68...	Fold1	1	<tibble [2,743 × ...	<tibble [686 × ...	<fit[+...
##	2 <split [2.7K/68...	Fold1	2	<tibble [2,743 × ...	<tibble [686 × ...	<fit[+...
##	3 <split [2.7K/68...	Fold1	3	<tibble [2,743 × ...	<tibble [686 × ...	<fit[+...
##	4 <split [2.7K/68...	Fold1	4	<tibble [2,743 × ...	<tibble [686 × ...	<fit[+...
##	5 <split [2.7K/68...	Fold1	5	<tibble [2,743 × ...	<tibble [686 × ...	<fit[+...
##	6 <split [2.7K/68...	Fold2	1	<tibble [2,743 × ...	<tibble [686 × ...	<fit[+...
##	7 <split [2.7K/68...	Fold2	2	<tibble [2,743 × ...	<tibble [686 × ...	<fit[+...
##	8 <split [2.7K/68...	Fold2	3	<tibble [2,743 × ...	<tibble [686 × ...	<fit[+...
##	9 <split [2.7K/68...	Fold2	4	<tibble [2,743 × ...	<tibble [686 × ...	<fit[+...
##	10 <split [2.7K/68...	Fold2	5	<tibble [2,743 × ...	<tibble [686 × ...	<fit[+...
##	# ... with 15 more rows					

## Model performance: post-grid search

```
rf_model_tun <- rf_tun %>%  
  mutate(pred=map2(.x=rf, .y=validate, ~predict(.x, .y) %>%  
    as_tibble %>%  
    set_names("pred"))  
    , mae=map2_dbl(.x=pred, .y=validate  
      , ~measures::MAE(truth = .y$quality  
        , response = .x$pred)  
    )  
    , mse=map2_dbl(.x=pred, .y=validate  
      , ~measures::MSE(truth = .y$quality  
        , response = .x$pred)  
    )  
  ))
```

# Model performance: post-grid search

```
rf_model_tun %>%  
  dplyr::select(mtry, id, mae, mse) %>%  
  pivot_longer(cols = c(mae, mse)  
               , names_to = "measure") %>%  
  ggplot(aes(as.factor(mtry), value)) +  
  geom_boxplot() +  
  geom_jitter(aes(col=id)) +  
  facet_grid(~measure, scales = "free")
```



## Regression: final model

```
best_model <- rand_forest(mtry = 5, trees = 200) %>%  
  set_engine("ranger", importance="impurity") %>%  
  fit(quality~.-quality_b, data=train_data)  
test_pred <- predict(best_model, test_data) %>%  
  as_tibble %>%  
  set_names("pred")  
  
mae_rf <- measures::MAE(truth = test_data$quality, response = test_pred$pred)  
mse_rf <- measures::MSE(truth = test_data$quality, response = test_pred$pred)
```


# Ridge

## Tweak data

```
X <- train_data %>%  
  dplyr::select(-quality, -quality_b) %>% as.matrix()  
y <- train_data %>%  
  dplyr::select(quality) %>% as.matrix()  
X_test <- test_data %>%  
  dplyr::select(-quality, -quality_b) %>% as.matrix()  
y_test <- test_data %>%  
  dplyr::select(quality) %>% as.matrix()
```

## Model

```
ridge_cv <- cv.glmnet(X, y, alpha = 0, standardize = T, nfolds
```




# Ridge

## Tweak model


```
ridge_lambda <- 10^seq(-3, 1, length.out = 100)
ridge_cv_a1 <- cv.glmnet(X, y, alpha = 0, lambda = ridge_lambda,
                        , standardize = T, nfolds = 10)
```

```
ridge_best <- glmnet(X, y, alpha = 0, lambda = ridge_cv_a1$lambda.1se,
                    , standardize = T)
ridge_pred <- predict(ridge_best, X_test)
mae_ridge <- measures::MAE(truth = test_data$quality, response = ridge_pred)
```



# Ridge


```
ridge_shrink <- glmnet(X, y, alpha = 0  
                      , lambda = ridge_lambda  
                      , standardize = F)
```






# Lasso

```
lasso_cv <- cv.glmnet(X, y, alpha = 1
                     , standardize = TRUE, nfolds = 10)
lasso_best <- glmnet(X, y, alpha = 1
                    , lambda = lasso_cv$lambda.min
                    , standardize = TRUE)
lasso_pred <- predict(lasso_best, X_test)
mae_lasso <- measures::MAE(truth = test_data$quality
                          , response = lasso_pred)
```




# Lasso shrinkage

```
lasso_lambda <- 10^seq(-3, 1, length.out = 100)
lasso_shrink <- glmnet(X, y, alpha = 1
  , lambda = lasso_lambda
  , standardize = F)
```



# Comparison

```
data.frame(linear=mae_lm, random_forest=mae_rf
           , ridge=mae_ridge, lasso=mae_lasso) %>%
  t() %>%
  as.data.frame() %>%
  tibble::rownames_to_column("MAE") %>%
  as_tibble() %>%
  ggplot(aes(reorder(MAE, desc(V1)), V1))+
  geom_col()+
  labs(x="Algorithm", y="Mean Absolute Error")
```



# References

- Géron, A. (2017). Hands-On Machine Learning with Scikit-Learn & TensorFlow.
- Hastie, T., Tibshirani, R., & Friedman, J. (2009). The Elements of Statistical Learning. Springer, 27(2), 83–85.
- James, G., Witten, D., Hastie, T., & Tibshirani, R. (2013). An Introduction to Statistical Learning. Springer. New York.
- Müller, A. C., & Guido, S. (2017). Introduction to machine learning with Python. O'Reilly.